

COMP3141 Assignment 2



Haskell-Augmented Regular Expressions (H.A.R.E.)

Version 1.3

Liam O'Connor

Term 2, 2019

Marking Total of 20 marks (20% of practical component)

Due Date Wednesday, 7th August, 2019, 09:00

Late Penalty The maximum available mark is reduced by 10% if the assignment is one day late, by 25% if it is 2 days late and by 50% if it is 3 days late. Assignments that are late 4 days or more will be awarded zero marks. So if your assignment is worth 88% and you submit it one day late you still get 88%, but if you submit it two days late you get 75%, three days late 50%, and four days late zero.

Submission Instructions The assignment can be submitted using the ‘give’ system.

To submit from a CSE terminal, type:

```
$ give cs3141 Hare Hare.hs
```

Overview

In this assignment, you will refactor a text-matching *regular expressions* engine to use a more strongly-typed interface. These stronger types not only reduce the possibility for bugs, but also improve the readability and clarity of the code base. On top of the strongly-typed version, you will implement a number of additional convenience combinators for working with regular expressions. This assignment will give you experience using Haskell type system extensions (specifically GADTs), experience programming with *monads* and *applicative functors*, and introduce you to concepts such as *alternatives*.

Provided Code

The provided code consists of a number of modules:

`Hare.hs` contains the stubs of the code you should implement.

`Untyped.hs` contains an untyped, unextended version of the regex engine.

`HareMonad.hs` contains the monadic type used to write the matching algorithm.

`Tests.hs` contains the `main` function for running tests.

`Tests/Support.hs` contains the support code for running tests.

`Tests/UnitTests.hs` contains properties for the basic regular expressions.

`Tests/Transcript.hs` contains some acceptance tests for your combinators, for analysing a UNSW transcript.

`Tests/Examples.hs` contains all examples from this spec, as unit tests.

Note: The **only** file you can submit is '`Hare.hs`', so make sure your submission compiles with the *original* versions of all other files.

Regular Expressions

Regular expressions are a popular way to describe and manipulate patterns of strings. Often they are used to search for a substring that matches a particular pattern, and then extracting information from that substring.

Many languages include regular expressions as a built in feature or in their standard library. If you have done COMP2041, you will have been exposed to several such languages. Unfortunately, Haskell is not such a language. To remedy this, Helen the Hare started to design H.A.R.E, a regular expressions library for Haskell.

She started by encoding the basic building blocks of regular expressions as a Haskell data type.

```
data RE = Empty           -- Matches the empty string
        | Fail            -- Matches no strings
        | Char [Char]     -- Matches a single character from the list
        | Seq RE RE       -- Matches the two expressions in sequence
        | Choose RE RE    -- Matches either the first, or the second
        | Star RE         -- Matches the expression zero or more times
```

Note: If you have seen regular expressions before, you may be used to more features being available than the ones above. In general, those features can be translated into the features given above.

Some simple examples of these regular expressions are given in the table below.

Regular Expression	Matches
<code>Star (Char ['0'..'9']) `Seq` Char ['0'..'9']</code>	<code>"1234","039","0",...</code>
<code>Star (Char ['a']) `Choose` Star (Char ['b'])</code>	<code>"","a","b","aa","bb"...</code>
<code>Fail</code>	<code>-- nothing</code>
<code>Empty</code>	<code>""</code>

In order to define a matching algorithm for regular expressions, Helen defined a type called `Hare` (found in `HareMonad.hs`), which is similar to `State String` except that the return type is wrapped in a monad `f`:

```
newtype Hare f a = Hare { runHare :: String -> f (String, a) }
```

The `hare` function can be used to “run” a `Hare` computation on a string:

```
hare :: Functor f => Hare f a -> String -> f a
hare a s = fmap snd (runHare a s)
```

Helen has parameterised the `Hare` type by `f` so that it could be instantiated to `Maybe` if the user only wishes to find the *first* match for a particular regular expression, or to the list type constructor `[]` if the user wishes to find *all possible* matches of the regular expression. This requires us to make use of the `Alternative` type class, which has two built in methods:

```
class Applicative f => Alternative f where
    empty :: f a
    (<|>) :: f a -> f a -> f a
```

The `empty` method denotes *failure*. For the list instance, it is the empty list `[]`, whereas for the `Maybe` instance it is `Nothing`. The `(<|>)` method denotes *alternation*. For the list instance it is the same as concatenation `(++)`, whereas for the `Maybe` instance, it is a left-biased choice, like so:

```
instance Alternative Maybe where
    empty = Nothing
    Just a <|> b = Just a
    Nothing <|> b = b
```

To avoid confusing `empty` with the `Empty` constructor for regular expressions, Helen also defined an alias for `empty`, called `failure`:

```
failure :: (Alternative f, Monad f) => Hare f a
failure = empty
```

The `readCharacter` action is also of note, which removes a character from the `String` state and returns it. If there are no characters left in the string, it will fail:

```
readCharacter :: (Alternative f, Monad f) => Hare f Char
readCharacter = Hare $ \s -> case s of
  [] -> empty
  (x:xs) -> pure (xs,x)
```

Armed with just `readCharacter` and the `Monad`, `Applicative` and `Alternative` instances for `Hare`, Helen defines the matching algorithm, in a function called `match` (found in `Untyped.hs`):

```
match :: (Monad f, Alternative f) => RE -> Hare f Results
```

Here the `Results` type describes the string that was matched by each part of the regular expression:

```
data Results = None
             | Character Char
             | Tuple Results Results
             | Repetition [Results]
```

Seeing as the `Empty` expression matches any string, Helen simply returns `None` in that case:

```
match Empty = pure None
```

Conversely, because `Fail` matches no strings at all, Helen always calls `failure` if attempting to match this expression:

```
match Fail = failure
```

To define the case for `Char`, Helen uses the built-in function `guard`, which will fail if the given condition evaluates to false:

```
match (Char cs) = do
  x <- readCharacter
  guard (x `elem` cs)
  pure (Character x)
```

To define the case for `Seq`, Helen makes use of the `Monad` instance, first matching on the first expression `a`, then the second `b`:

```
match (Seq a b) = do
  ra <- match a
  rb <- match b
  pure (Tuple ra rb)
```

Because there is no dependencies between `ra` and `rb`, this could also have been defined using the `Applicative` instance, as follows:

```
match (Seq a b) = Tuple <$> match a <*> match b
```

For the case for `Choose`, Helen makes use of the `Alternative (<|>)` method:

```
match (Choose a b) =
  match a
  <|> match b
```

For `Star`, Helen observed the following bijection¹:

$$\text{Star } r \quad === \quad (r \text{ `Seq` Star } r) \text{ `Choose` Empty}$$

For this reason, Helen implements `Star a` by first matching the expression `a`, then matching on `Star a` recursively. If either match fails, it falls back to returning the empty match:

```
match (Star a) =
  addFront <$> match a <*> match (Star a)
  <|> pure (Repetition [])
where
  addFront x (Repetition xs) = Repetition (x:xs)
  addFront _ _ = error "(should be) impossible!"
```

Note that Helen defined `Star` to choose the *longer alternative* first so that if used with `Maybe`, the `match` function will return the longest match rather than the empty string.

Having defined `match`, Helen can now define the regex matching operator (`=~`), which searches through a string for a match to the expression:

```
(=~) :: (Monad f, Alternative f) => String -> RE -> f Results
str =~ re = hare matchAnywhere str
  where matchAnywhere = match re <|> (readCharacter >> matchAnywhere)
```

If we look at some examples, we can see that invoking it with the `Maybe` monad returns just the first result, whereas the list version returns all results:

```
*> "COMP3141" =~ (Char ['0'..'9']) :: Maybe Results
Just (Character '3')
*> "COMP3141" =~ (Char ['0'..'9']) :: [Results]
[Character '3',Character '1',Character '4',Character '1']
```

¹While they return different results, there is a bijection between them

A Typed Implementation (12 Marks)

Helen's untyped implementation *works*, but it's not very clean. For example, we *know* that the `Results` returned by the regular expression

```
Char ['0'..'9'] `Seq` Star (Char ['0'..'9'])
```

will be of the format:

```
Tuple (Character c) (Repetition [Character x, Character y, ...])
```

but any function designed to process the `Results` from this regular expression would have to be *partial*, because we did not encode the format of the results returned into the type system.

Instead of having a `Results` type, your first task will be to define a *type-indexed* version of the RE type and associated `match` function, in `Hare.hs`.

```
data RE :: * -> * where
  Empty :: RE ()
  Fail  :: RE a
  -- Your constructors here
```

Each constructor has an additional type argument which specifies the type of the result returned from each expression. We have defined the first two constructors for you. You will have to find appropriate definitions for the rest of them.

For implementing `match`, feel free to look at Helen's original implementation in the `Untyped.hs` module as a guide. The test cases provided may also be instructive in getting your implementation type-checking and correct.

Marking Criteria	
Marks	Description
2	Empty type correct; <code>match</code> implementation correct.
2	Fail type correct; <code>match</code> implementation correct.
2	Char type correct; <code>match</code> implementation correct.
2	Seq type correct; <code>match</code> implementation correct.
2	Choose type correct; <code>match</code> implementation correct.
2	Star type correct; <code>match</code> implementation correct.
12	Total

Adding Actions (1 Marks)

Next, we will add an additional constructor to `RE` that does not change the set of strings matched, but allows us to run an arbitrary transformation on the results returned from the match:

```
Action :: (a -> b) -> RE a -> RE b
```

Your task in this part is to implement the `match` function for this constructor. You may find the `Functor` instance for `Hare` useful.

Examples

```
*> "***" =~ Action length (Star (Char ['*'])) :: [Int]
[3,2,1,0,2,1,0,1,0,0]

*> "AXax" =~ Action isUpper (Char ['a','A']) :: [Bool]
[True, False]

*> let f (x,y) = [x,y]
*> let atoz = Char ['a'..'z']
*> "ab01cd20" =~ Action f (atoz `Seq` atoz) :: [String]
["ab","cd"]
```

Marking Criteria	
Marks	Description
1	Action clause in <code>match</code> correct.
1	Total

Utility Combinators (7 Marks)

Using our newly minted `Action` constructor, you must now define a series of combinator functions that allow us to define a lot of the features people normally expect to find in a regular expressions system.

`cons` function

The `cons` function matches a regular expression for an element of type `a`, followed by a regular expression for a list `[a]`, and returns that element prepended to the list.

```
cons :: RE a -> RE [a] -> RE [a]
```

Examples:

```
*> "10100" =~ cons (Char ['1']) (Star (Char ['0'])) :: [String]
["10", "1", "100", "10", "1"]

*> "10100" =~ cons (Char ['1']) (Action (const []) Empty) :: [String]
["1", "1"]
```

`plus` function

The `plus` function is like `Star`, but requires the expression given to occur *at least once*. Thus, this function can be defined using `Action`, `Seq` and `Star`.

```
plus :: RE a -> RE [a]
```

Examples:

```
*> "10100" =~ plus (Char ['0']) :: [String]
["0", "00", "0", "0"]

*> let atoz = Char ['a'..'z']
*> let digits = Char ['0'..'9']
*> "ab1c3" =~ plus (atoz `Seq` digits) :: [(Char,Char)]
[(('b', '1'), ('c', '3')), (('b', '1')), (('c', '3'))]
```

`string` function

The `string` function produces a regex that only matches the given string, and returns it. You should be able to implement it using `Char`, `cons`, `Empty` and `Action`.

```
string :: String -> RE String
```

Examples:

```
*> let comp3141 = string "COMP3141"
*> "My favourite subject is COMP3141" =~ comp3141 :: Maybe String
Just "COMP3141"
*> "My favourite subject is MATH1141" =~ comp3141 :: Maybe String
Nothing
```


choose function

The `choose` function is like the `Choose` constructor, but generalised to lists. That is, `choose [a, b]` is equivalent to a `'Choose'` `b`. What should `choose []` be?

```
choose :: [RE a] -> RE a
```

Examples:

```
*> let re = choose [string "COMP", string "MATH", string "PHYS"]
*> "COMP3141, MATH1081, PHYS1121, COMP3121" =~ re :: [String]
["COMP", "MATH", "PHYS", "COMP"]
```

```
*> "abc" =~ choose [] :: Maybe String
Nothing
```

option function

The `option` function matches the given expression zero or one times. In other words, if the expression matches, it returns `Just` that match, otherwise `Nothing`. You can implement this combinator with `Action`, `Choose` and `Empty`.

```
option :: RE a -> RE (Maybe a)
```

Examples:

```
*> let digits = Char ['0'..'9']
*> let sign = Action (fromMaybe '+') (option (Char ['-']))
*> "-30 5 3" =~ (sign `cons` plus digits) :: [String]
["-30", "-3", "+30", "+3", "+0", "+5", "+3"]
```

```
*> "foo" =~ option (Char ['a']) :: [Maybe Char]
[Nothing, Nothing, Nothing, Nothing]
```

Note: As with `Star`, prefer *longer* matches over shorter ones, so that the results appear as in the above example.

rpt function

The `rpt` combinator allows a regular expression to be repeated a fixed number of times. The expression `rpt n x` is equivalent to repeating `x` in sequence `n` times, returning the results in a list.

```
rpt :: Int -> RE a -> RE [a]
```

Examples:

```

*> let digits = Char ['0'..'9']
*> let programs = choose [string "COMP", string "PHYS", string "MATH"]
*> let courseCode = programs `Seq` rpt 4 digits
*> "COMP3141, MATH1081, and PHYS1121" =~ courseCode :: [(String,String)]
[("COMP", "3141"), ("MATH", "1081"), ("PHYS", "1121")]

*> "foo" =~ rpt 0 (Char ['a']) :: Maybe [Char]
Just ""

```

rptRange function

Lastly, the `rptRange` function takes a range of numbers (x, y) . You may assume that $x \leq y$. It will match the given regex at least x times and at most y times.

```
rptRange :: (Int, Int) -> RE a -> RE [a]
```

Examples:

```

*> "1234" =~ rptRange (2,4) (Char ['0'..'9']) :: [String]
["1234", "123", "12", "234", "23", "34"]

*> "1234" =~ rptRange (3,3) (Char ['0'..'9']) :: [String]
["123", "234"]

```

Note: As with `Star` and `option`, prefer longer matches to shorter ones.

Marking Criteria	
Marks	Description
1	<code>cons</code> implementation correct.
1	<code>plus</code> implementation correct.
1	<code>string</code> implementation correct.
1	<code>choose</code> implementation correct.
1	<code>option</code> implementation correct.
1	<code>rpt</code> implementation correct.
1	<code>rptRange</code> implementation correct.
7	Total

Compiling and Building

In addition to the standard library, this project depends on the `QuickCheck` and `HUnit` testing libraries and the test framework called `tasty`. For CSE machines, we have already configured a package database on the course account that should build the assignment without difficulty using the standard Haskell build tool `cabal`. For students using their own computer, we instead recommend the alternative build tool `stack`, available from the Haskell Stack website at <https://www.haskellstack.org>. We have provided a `stack` configuration file that fixes the versions of each of our dependencies to known-working ones. If you use `stack` to set up your toolchain, you should have minimal compatibility difficulties regardless of the platform you are using. *If you are using versions of GHC or Haskell build tools supplied by your Linux distribution, these are commonly out of date or incorrectly configured. We cannot provide support for these distributions.*

Detailed, assignment-specific instructions for each build tool are presented below.

On CSE Machines

Enter a COMP3141 subshell by typing `3141` into a CSE terminal:

```
$ 3141
newclass starting new subshell for class COMP3141...
```

From there, if you navigate to the directory containing the assignment code, you can build the assignment by typing:

```
$ cabal build
```

To run the program from `Tests.hs`, which contains all the unit, acceptance, and property tests, type

```
$ ./dist/build/HareTests/HareTests
```

To start a `ghci` session, type:

```
$ cabal repl
```

Lastly, if for whatever reason you want to remove all build artefacts, type:

```
$ cabal clean
```

For stack users

Firstly, ensure that GHC has been setup for this project by typing, in the directory that contains the assignment code:

```
$ stack setup
```

If `stack` reports that it has already set up GHC, you should be able to build the assignment with:

```
$ stack build
```

This `build` command will, on first run, download and build the library dependencies as well, so be sure to have an internet connection active. To run the program from `Tests.hs`, which contains all the unit, acceptance, and property tests, type

```
$ stack exec HareTests
```

To start a `ghci` session, type:

```
$ stack repl
```

Lastly, if for whatever reason you want to remove all build artefacts, type:

```
$ stack clean
```

Marking and Testing

Testing

A number of different test suites are provided:

- The `Tests/Transcript.hs` file contains some high-level *acceptance* tests that extract various bits of information from an anonymised UNSW student transcript, located in `Tests/transcript.txt`.
- The `Tests/UnitTests.hs` file contains some QuickCheck properties that, while not complete specifications, will help you to gain some assurance for your implementation of `match` for each constructor.
- The `Tests/Examples.hs` file contains each example presented in this spec document as a unit test.
- The `Tests.hs` file contains a `tasty` test suite containing all of the above tests.

Marking

All marks for this assignment are awarded based on *automatic marking scripts*, which are comprised of QuickCheck properties, acceptance tests, and unit tests. Marks are not awarded subjectively, and are allocated according to the criteria presented in each section.

Barring exceptional circumstances, the marks awarded by the automatic marking script are *final*. For this reason, please make sure that your submission compiles and runs correctly on CSE machines. We will use similar machines to mark your assignment.

A dry-run script that runs the tests provided in the assignment code will be provided. When you submit the assignment, please make sure the script does not report any problems.

Late Submissions

Unless otherwise stated if you wish to submit an assignment late, you may do so, but a late penalty reducing the maximum available mark applies to every late assignment. The maximum available mark is reduced by 10% if the assignment is one day late, by 25% if it is 2 days late and by 50% if it is 3 days late. Assignments that are late 4 days or more will be awarded zero marks. So if your assignment is worth 88% and you submit it one day late you still get 88%, but if you submit it two days late you get 75%, three days late 50%, and four days late zero.

Extensions

Assignment extensions are only awarded for serious and unforeseeable events. Having the flu for a few days, deleting your assignment by mistake, going on holiday, work commitments, etc do not qualify. Therefore aim to complete your assignments well before the due date in case of last minute illness, and make regular backups of your work.

Plagiarism

Many students do not appear to understand what is regarded as plagiarism. This is no defense. Before submitting any work you should read and understand the UNSW plagiarism policy <https://student.unsw.edu.au/plagiarism>.

All work submitted for assessment must be entirely your own work. We regard unacknowledged copying of material, in whole or part, as an extremely serious offence. In this course submission of any work derived from another person, or solely or jointly written by and or with someone else, without clear and explicit acknowledgement, will be severely punished and may result in automatic failure for the course and a mark of zero for the course. Note this includes including unreferenced work from books, the internet, etc.

Do not provide or show your assessable work to any other person. Allowing another student to copy from you will, at the very least, result in zero for that assessment. If you knowingly provide or show your assessment work to another person for any reason, and work derived from it is subsequently submitted you will be penalized, even if the work was submitted without your knowledge or consent. This will apply even if your work is submitted by a third party unknown to you. You should keep your work private until submissions have closed.

If you are unsure about whether certain activities would constitute plagiarism ask us before engaging in them!